

App Note - VC5 MMU

The VC5 includes an MMU on 7271x0 / 7268x0 / 7260A0 and 7278A0. This should lead to improved performance in a CMA based system (such as Android), is immune to fragmentation and increases the memory available to Linux for applications. It also provides cross process memory protection and coherency required for upcoming standards such as Vulkan.

NEXUS_MEMC0_GRAPHICS_HEAP heaps on these platforms have been shrunk to provide only space for interoperability surfaces, such as output framebuffers. Currently they are 64MB, so should be suitable to run any of the BSEAV/lib/gpu/applications samples in a non MMU configuration.

Requirements

- 1) BOLT image v1.25 or later

The MMU driver is based on the standard Linux DRM framework (https://en.wikipedia.org/wiki/Direct_Rendering_Manager) and requires a device tree to operate. **BOLT versions from v1.25 will auto configure.**

Prior versions can be used with the following commands to add entries at BOLT> command prompt

```
dt add node rdb gpu
dt add prop rdb/gpu compatible s 'brcm,v3d-v3.3.0.0'
```

Version number is important, so contact the 3d team for your specific platform.

- 2) Linux kernel 4.1-1.7 or later

4.1-1.7 kernel shipped with CONFIG_DRM enabled as default.

Prior versions can be used by rebuilding the kernel with the requisite option above. Our Linux kernel ships with .config copied to config, which various parts of the Nexus build rely on. Nexus also needs to be rebuilt against the modified version as the module will fail to install. In some versions of the kernel the CMA driver is initialized with all the memory allocated, it is therefore necessary to use 'cmatool resetall' at boot to free the CMA allocation. **However later 4.1 kernels (e.g. 4.1-1.12) and the 4.9 kernels do not require this step.**

- 3) Boot line

The shared memory between CPU and GPU is determined via the boot parameters. Nexus would ordinarily suggest a config, such as 'vmalloc=374m bmem=696m@1352m'. As an example we may wish to allocate a shared region of 512MB before the BMEM region. drm.debug is optional and is suggested the first time you run, giving additional verbosity on the kernel log.

```
boot stbbri-fs-1:/41-1.6/vmlinuz-initrd-arm 'vmalloc=374m  
bmem=696m@1352m brcm_cma=512m@820m drm.debug=0x2'
```

Running

No additional changes are required to run your application. Some new console output is visible and additional debug tools are available.

Is the MMU being used?

When starting an application, you should see on the console.

DRM: Trying to open: /dev/dri/card0

```
*** 00:00:03.342 nexus_statistics_callback: TaskCallback callback 0x79b38:0xb5e99cf8 from  
modules/hdmi_output/src/nexus_hdmi_output.c:457 202334 usec
```

```
*** 00:00:03.342 nexus_statistics_module: display[Default]  
nexus/core/syncthunk/nexus_display_thunks.c:26 201727 usec
```

```
*** 00:00:03.342 nexus_statistics_api: NEXUS_Display_GetSettings[display:Default] 201816  
usec
```

DRM: brcmv3d Version 1.1.0 20161207, Broadcom V3D GEM provider

DRM: No Nexus secure graphics heap available

If you observe the message - “Unable to open device - defaulting back to Nexus”, then the device is in legacy mode.

It should be working but is not

Below are some simple steps to triage a system if you believe the MMU is not being used.

- 1) Check that the kernel module is built and present in NEXUS_BIN_DIR
For example obj.97271/nexus/bin/brcmv3d.ko
- 2) Check the module can be installed on the STB
insmod brcmv3d.ko
- 3) Check the node is present on the STB
ls /dev/dri/card0

If you get any error messages with any of these steps, report them to the 3d team.

Debug capabilities

The MMU driver exposes its config via the debug fs.

Checking what is currently using the driver.

```
# cat /sys/kernel/debug/dri/0/clients
      command  pid dev master a   uid   magic
      cube    1448  0   y   y    0     0
```

Total allocations made on a client

```
# cat /sys/kernel/debug/dri/0/<PID>-<MAGIC>/client
```

Here, a single instance of cube is running.

Checking allocations made by a client

```
# cat /sys/kernel/debug/dri/0/<PID>-<MAGIC>/objs
```

handle	size	HW	Virtual	RO	WC	EXT	DESC
1	67174400	0x00010000		n	n	y	NEXUS Offscreen Heap
2	1048576	0x04020000		n	n	n	
3	65536	0x04120000		y	n	n	
5	65536	0x04290000		n	n	n	
6	65536	0x042a0000		n	n	n	
7	65536	0x042b0000		y	n	n	
8	65536	0x042c0000		y	n	n	
9	65536	0x042d0000		y	n	n	
10	65536	0x042e0000		y	n	n	
11	65536	0x042f0000		y	n	n	
12	1441792	0x04300000		n	n	y	
13	65536	0x04460000		y	n	n	
14	65536	0x04470000		y	n	n	
15	1441792	0x04480000		n	n	y	
16	65536	0x045e0000		y	n	n	
17	65536	0x045f0000		y	n	n	

The Nexus offscreen heap is always mapped into the address space so the device can write to the interop surfaces.

Checking current 2MB pages

The DRM driver uses 2MB blocks, and then breaks them internally to a smaller granularity.

```
# cat /sys/kernel/debug/dri/0/<PID>-<MAGIC>/cma
CMA Block Physical:      64K Page Allocation Mask
=====
0x000000002f000000: *****-----
```

In this instance, cube is running using a single 2MB CMA block, with four 64kb pages remaining.